
Instruction Manual for the Application

Md Arafat Hossain Khan

arafat.math@gmail.com

December, 2020

Contents

| | | |
|----------|---|-----------|
| 1 | Instruction to run and test the application on AWS | 1 |
| 2 | Development | 6 |
| 2.1 | Development Stage on Databricks | 6 |
| 2.2 | Development Stage on IntelliJ IDEA | 6 |
| 2.2.1 | Versions | 6 |
| 2.2.2 | Dependencies | 6 |
| 2.2.3 | build.sbt file | 7 |
| 2.2.4 | Scala Class | 8 |
| 2.3 | Creating jar file | 8 |
| 3 | Trancated Sample Output | 10 |

Chapter 1

Instruction to run and test the application on AWS

Step 1: Data Collection

Go to https://www.transtats.bts.gov/DL_SelectFields.asp?Table_ID=236. Select 'Filter Year' and 'Filter Period' from dropdown menu. Then from the 'Field Name' column, select only three columns ORIGIN, DEST and REPORTING_AIRLINE. Make sure that all other fields are unchecked. Click on 'Download' button. A zipped folder will be downloaded. Extract the csv file from this zipped folder.

Step 2: Log in to AWS

Go to <https://aws.amazon.com/>, log in to your account and go to 'AWS Management Console'.

Step 3: Create Key Pair

From EC2 Dashboard, create key pair.

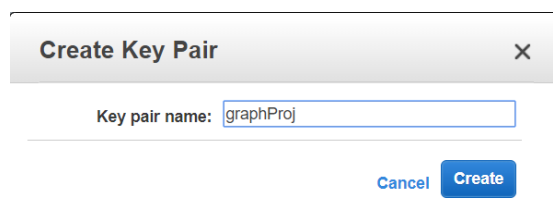


Figure 1.1: Create Key Pair.

Save the 'graphProj.pem' file for later use if needed. But in this testing we will not use it locally.

Step 4: Create Cluster

Now follow the following setup to create a cluster from 'Amazon EMR'.

General Configuration

Cluster name **graphProject**

☒ Logging

S3 folder **s3://aws-logs-761776884662-us-east-1/elasticmapred**

Launch mode ☒ Cluster ☐ Step execution

Software configuration

Release **emr-5.23.0**

Applications

- ☐ Core Hadoop: Hadoop 2.8.5 with Ganglia 3.7.2, Hive 2.3.4, Hue 4.3.0, Mahout 0.13.0, Pig 0.17.0, and Tez 0.9.1
- ☐ HBase: HBase 1.4.9 with Ganglia 3.7.2, Hadoop 2.8.5, Hive 2.3.4, Hue 4.3.0, Phoenix 4.14.1, and ZooKeeper 3.4.13
- ☐ Presto: Presto 0.214 with Hadoop 2.8.5 HDFS and Hive 2.3.4 Metastore
- ☒ **Spark: Spark 2.4.0 on Hadoop 2.8.5 YARN with Ganglia 3.7.2 and Zeppelin 0.8.1**
- ☐ Use AWS Glue Data Catalog for table metadata

Hardware configuration

Instance type **m3.xlarge**

Number of instances **3** (1 master and 2 core nodes)

Security and access

EC2 key pair **graphProj**

Permissions ☒ Default ☐ Custom

Use default IAM roles. If roles are not present, they will be automatically created for you with managed policies for automatic policy updates.

EMR role **EMR_DefaultRole**

EC2 instance profile **EMR_EC2_DefaultRole**

[Cancel](#) [Create cluster](#)

Figure 1.2: Create Cluster.

At the beginning the window will show the following,

Cluster: graphProject **Starting**

Summary Application history Monitoring Hardware Configurations Events Steps Bootstrap actions

Connections: --

Master public DNS: --

Tags: -- [View All / Edit](#)

Summary

ID: j-2C4FOGLCTR2PQ
Creation date: 2019-04-25 04:05 (UTC-5)
Elapsed time: 1 second
Auto-terminate: No
Termination protection: Off [Change](#)

Configuration details

Release label: emr-5.23.0
Hadoop distribution: Amazon
Applications: Ganglia 3.7.2, Spark 2.4.0, Zeppelin 0.8.1
Log URI: s3://aws-logs-761776884662-us-east-1/elasticmapreduce/

EMRFS consistent view: Disabled
Custom AMI ID: --

Network and hardware

Availability zone: --
Subnet ID: **subnet-d2e492b5**
Master: Provisioning 1 m3.xlarge
Core: Provisioning 2 m3.xlarge

Security and access

Key name: graphProj
EC2 instance profile: EMR_EC2_DefaultRole
EMR role: EMR_DefaultRole
Visible to all users: All [Change](#)
[This feature will be deprecated soon](#)

Security groups for Master: --
Security groups for Core & Task: --

Figure 1.3: Creating Cluster.

When the cluster will be ready, the window will show the following,




Figure 1.4: Cluster Created.

Step 5: S3 Bucket

Now, upload the csv file you downloaded and the jar file I uploaded on eLearning.

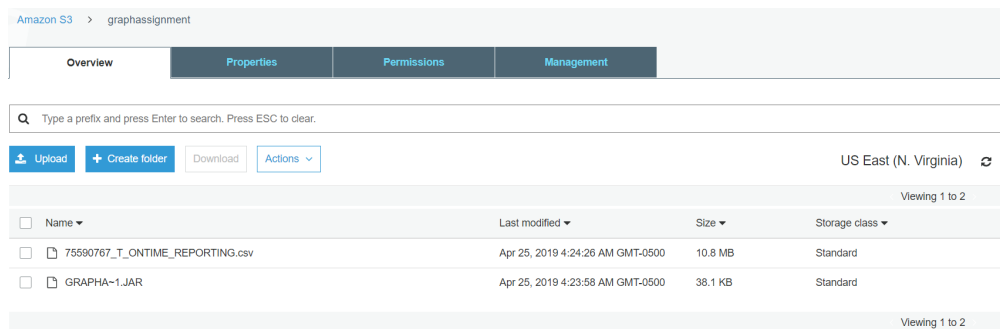


Figure 1.5: S3 Bucket and File Upload.

Step 6: Add Step for Execution and Provide Input Output File Path as Argument

Now, as the cluster is running and the required files are uploaded we are ready to add a step to execute the application and providing the data file path and output path. Following table and the figure demonstrates the setup,

```

1 --class graphAssignmentArafat --packages graphframes:graphframes
  :0.7.0 --spark2.4-s_2.11
2
3 s3://graphassignment/graphassignmenttarafat_2.11-0.1.jar
4
5
6 s3://graphassignment/75590767_T_ONTIME_REPORTING.csv
7 s3://graphassignment/output

```

Add step [X]

Step type Spark application

Name Graph Spark application

Deploy mode Cluster

Spark-submit options `--class graphAssignmentArafat --packages graphframes:graphframes:0.7.0-spark2.4-s_2.11`

Application location s3://graphassignment/graphassignmentararat_2.11-0.1

Arguments `s3://graphassignment/75590767_T_ONTIME_REPORTING.csv`
`s3://graphassignment/output`

Action on failure Continue

Cancel **Add**

Figure 1.6: Add Step.

After the step is added the status of the execution will start. This program takes around 1.5 hours to execute.

Step 7: Check Output

You will see a folder named 'output' in your bucket on Amazon S3 once the program is completed. There are multiple files in the folder. These files can be downloaded to read the output. In my case there were

Now, finally, we are ready to see the output. Here, the output files can be opened using Notepad++ or Wordpad application. It is important to note that the output files should not be opened using Notepad. Because the formatting will better be demonstrated in Notepad++ or Wordpad. Notepad may not be able to handle Unix newlines. Chapter 3 shows the truncated sample output.














| Viewing 1 to 13 | | | |
|--|-----------------------------------|---------|-----------------|
| <input type="checkbox"/> Name ▼ | Last modified ▼ | Size ▼ | Storage class ▼ |
| <input type="checkbox"/>  _SUCCESS | Apr 24, 2019 11:28:53 PM GMT-0500 | 0 B | Standard |
| <input type="checkbox"/>  part-00000 | Apr 24, 2019 11:28:52 PM GMT-0500 | 932.0 B | Standard |
| <input type="checkbox"/>  part-00001 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input type="checkbox"/>  part-00002 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input type="checkbox"/>  part-00003 | Apr 24, 2019 11:28:52 PM GMT-0500 | 932.0 B | Standard |
| <input type="checkbox"/>  part-00004 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input type="checkbox"/>  part-00005 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input type="checkbox"/>  part-00006 | Apr 24, 2019 11:28:52 PM GMT-0500 | 932.0 B | Standard |
| <input type="checkbox"/>  part-00007 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input type="checkbox"/>  part-00008 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input type="checkbox"/>  part-00009 | Apr 24, 2019 11:28:52 PM GMT-0500 | 932.0 B | Standard |
| <input type="checkbox"/>  part-00010 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |
| <input checked="" type="checkbox"/>  part-00011 | Apr 24, 2019 11:28:52 PM GMT-0500 | 934.0 B | Standard |

Figure 1.7: Output Folder.

Chapter 2

Development

2.1 Development Stage on Databricks

First I developed the code on Databricks. In this stage I explicitly coded everything and used the had coded path variables of the data files. The code is as follows,

```
1 import org.apache.spark.sql.functions._
2 import org.apache.spark.graphx._
3 import org.apache.spark.rdd.RDD
4 import org.graphframes._
5 val airport = spark.read.option("header", "true").option("inferSchema", "true").csv("/FileStore/tables/84936234_
  _T_ONTIME_REPORTING.csv")
6 val graphTable = airport.groupBy("ORIGIN", "DEST").count().orderBy(desc("count")).toDF("Origin", "Dest", "count")
7
8 -----I SKIPPED THE LONG MIDDLE PORTION OF THE CODE -----
9
10 val motifs = g.find("(A)-[]->(B); (B)-[]->(C); !(A)-[]->(C)")
11 var ABCList = motifs.select("A.name", "B.name", "C.name").toDF("A", "B", "C")
12 var Q6 = "\n\n6. Triplets [A, B, C] such that, at least one flight from A to B, and from B to C, but no flight from
  A to C.\n"
13 var resultABC = ABCList.rdd.map(_._1.toString())
14 Q6 += resultABC.collect().deep.mkString(",")
15 // ----- Final Answer ----- //
16 var Answer = Q1+Q2+Q3+Q4+Q5+Q6
17 println(Answer)
```

The output was as expected.

2.2 Development Stage on IntelliJ IDEA

2.2.1 Versions

First I created an empty project named `graphAssignmentArafat` where I used the following,

```
1 name := "\graphAssignmentArafat"
2 version := "0.1"
3 scalaVersion := "2.11.8"
4 val sparkVersion = "2.2.1"
```

2.2.2 Dependencies

For the project I also had to use the following dependencies,


```

1 import org.apache.spark.sql.functions._
2 import org.apache.spark.graphx._
3 import org.apache.spark.rdd.RDD
4 import org.apache.spark.sql.SQLContext
5 import org.graphframes._
6 import org.apache.spark.{SparkConf, SparkContext}

```

`org.apache.spark.sql.functions` has been imported for utilizing the functions available for DataFrame operations.

`org.apache.spark.graphx._` has been imported to utilize the GraphX library. This is a new component in Spark for graphs and graph-parallel computation. In addition, GraphX includes a growing collection of graph algorithms and builders to simplify graph analytics tasks.

A Resilient Distributed Dataset (RDD), the basic abstraction in Spark. That is why we need `org.apache.spark.rdd.RDD` for our spark application where RDD is in the core.

GraphFrames is a package for Apache Spark which provides DataFrame-based Graphs. Specially GraphFrames also support new algorithms:

1. Breadth-first search (BFS): Find shortest paths from one set of vertices to another
2. Motif finding: Search for structural patterns in a graph.

We will utilize Motif finding in our application. That is why we need `org.graphframes._` to be imported.

2.2.3 build.sbt file

I have created the following `build.sbt` file,

```

1 /*
2  * Author: Md Arafat Hossain Khan, PhD Candidate
3  * Department: Mathematics
4  */
5
6 name := "graphAssignmentArafat"
7
8 version := "0.1"
9
10 scalaVersion := "2.11.8"
11
12 // Extra Dependencies
13
14 val sparkVersion = "2.2.1"
15
16 // Refer to https://mvnrepository.com/
17
18 libraryDependencies += Seq(
19   "org.apache.spark" %% "spark-core" % sparkVersion,
20   "org.apache.spark" %% "spark-graphx" % sparkVersion,
21   "org.apache.spark" %% "spark-sql" % sparkVersion,
22   "org.apache.spark" %% "spark-mllib" % sparkVersion % "runtime",
23   "org.apache.spark" %% "spark-streaming" % sparkVersion % "provided",
24   "org.apache.spark" %% "spark-hive" % sparkVersion % "provided"
25 )
26
27 // Instruction for graphframes is available at https://spark-packages.org/package/graphframes/graphframes
28 resolvers += "Spark Packages Repo" at "http://dl.bintray.com/spark-packages/maven"
29 libraryDependencies += "graphframes" % "graphframes" % "0.7.0-spark2.4-s_2.11"
30
31 // https://mvnrepository.com/artifact/org.apache.spark/spark-graphx
32 // libraryDependencies += "org.apache.spark" %% "spark-graphx" % "2.2.1"
33 // https://mvnrepository.com/artifact/org.apache.spark/spark-sql
34 // libraryDependencies += "org.apache.spark" %% "spark-sql" % "2.2.1"
35 // https://mvnrepository.com/artifact/org.apache.spark/spark-mllib
36 // libraryDependencies += "org.apache.spark" %% "spark-mllib" % "2.2.1" % "runtime"
37 // https://mvnrepository.com/artifact/org.apache.spark/spark-streaming

```

```

38 // libraryDependencies += "org.apache.spark" %% "spark-streaming" % "2.2.1" % "provided"
39 // https://mvnrepository.com/artifact/org.apache.spark/spark-hive
40 // libraryDependencies += "org.apache.spark" %% "spark-hive" % "2.2.1" % "provided"

```

2.2.4 Scala Class

I created one class names `graphAssignmentArafat`. In this class I used the input and output file paths as arguments. So, while using the class on AWS or Databricks I kept the option for choosing the desired csv file as an input as well as its path. Also the output has been saved in a folder named `output`. The class looks like the following,

```

1  /*
2  Author: Md Arafat Hossain Khan, PhD Candidate
3  Department: Mathematics
4  */
5
6  import org.apache.spark.sql.functions._
7  import org.apache.spark.graphx._
8  import org.apache.spark.rdd.RDD
9  import org.apache.spark.sql.SQLContext
10 import org.apache.spark.sql.graphframes._
11 import org.apache.spark.{SparkConf, SparkContext}
12
13 object graphAssignmentArafat {
14   def main(args: Array[String]): Unit = {
15     if (args.length != 2) {
16       println("Usage: graphAssignmentArafat InputDir OutputDir")
17     }
18     // create Spark context with Spark configuration
19     val sc = new SparkContext(new SparkConf().setAppName("Spark graphAssignmentArafat"))
20     val sqlContext = new SQLContext(sc)
21     import sqlContext.implicits._
22
23     val Q0 = "This project was created by\n"+
24             "Md Arafat Hossain Khan\n"+
25             "PhD Candidate\n"+
26             "Department of Mathematics\nUniversity of Texas at Dallas"
27
28     /////////////// Code Start //////////////////////
29
30     -----I SKIPPED THE LONG MIDDLE PORTION OF THE CODE -----
31
32     /////////////// ----- Code End ----- //////////////////////
33
34     // Prepare output string
35     var answerRDD = sc.parallelize(Seq(Answer))
36
37     /////////////// Code End //////////////////////
38     answerRDD.map(r => r + "\n").saveAsTextFile(args(1))
39   }
40 }

```

2.3 Creating jar file

I used the command line to create the jar file. Following is the glimpse of the command line,

```

1 C:\Users\Md Arafat H Khan\Dropbox\Academics - University of Texas at Dallas\CS 6307.18s Introduction to Big Data
  Management and Analytics for non CS-Majors - Spring 2019\Assignment Solution\Assignment 03\
  graphAssignmentArafat>sbt
2 Java HotSpot(TM) 64-Bit Server VM warning: ignoring option MaxPermSize=256m; support was removed in 8.0
3 [info] Loading global plugins from C:\Users\Md Arafat H Khan\.sbt\1.0\plugins
4 [info] Loading project definition from C:\Users\Md Arafat H Khan\Dropbox\Academics - University of Texas at Dallas\
  CS 6307.18s Introduction to Big Data Management and Analytics for non CS-Majors - Spring 2019\Assignment
  Solution\Assignment 03\graphAssignmentArafat\project
5 [info] Loading settings from build.sbt ...
6 [info] Set current project to graphAssignmentArafat (in build file :C:/Users/Md%20Arafat%20H%20Khan/Dropbox/
  Academics%20-%20University%20of%20Texas%20at%20Dallas/CS%206307.18s%20Introduction%20to%20Big%20Data%20
  Management%20and%20Analytics%20for%20non%20CS-Majors%20-%20Spring%202019/Assignment%20Solution/Assignment
  %2003/graphAssignmentArafat/)
7 [info] sbt server started at local:sbt-server-cbb1816e2de0852a07f7
8 sbt:graphAssignmentArafat> compile

```

```
9 [info] Compiling 1 Scala source to C:\Users\Md Arafat H Khan\Dropbox\Academics - University of Texas at Dallas\CS
6307.18s Introduction to Big Data Management and Analytics for non CS-Majors - Spring 2019\Assignment Solution
\Assignment 03\graphAssignmentArafat\target\scala-2.11\classes ...
10 [warn] there was one deprecation warning; re-run with -deprecation for details
11 [warn] one warning found
12 [info] Done compiling.
13 [success] Total time: 27 s, completed Apr 25, 2019 3:30:09 AM
14 sbt:graphAssignmentArafat> package
15 [info] Packaging C:\Users\Md Arafat H Khan\Dropbox\Academics - University of Texas at Dallas\CS 6307.18s
Introduction to Big Data Management and Analytics for non CS-Majors - Spring 2019\Assignment Solution\
Assignment 03\graphAssignmentArafat\target\scala-2.11\graphassignmentarafat_2.11-0.1.jar ...
16 [info] Done packaging.
17 [success] Total time: 2 s, completed Apr 25, 2019 3:30:19 AM
18 sbt:graphAssignmentArafat>
```

`\graphAssignmentArafat\target\scala-2.11\` folder contains the jar file named `graphassignmentarafat_2.11-0.1.jar`.

Then follow the process of Chapter 1 to execute on AWS.

Chapter 3

Truncated Sample Output

The actual output is much bigger. I am presenting here the sample output read by Notepad++ after downloading from AWS.

```
1 1. Find the total number of airports (vertices) and the total flights
   connecting them.
2 Total number of airports :: 295
3 Total number of flights :: 464205
4
5 2. Find the in-degree and out-degree of each airport.
6 in-degree :: (ATL,151),(ORD,148),(DEN,124),(MSP,112),(DFW,106),(IAH
   ,103),(DTW,99),(PHX,89),(SLC,88),(EWR,85),(SFO,79),(LAX,79),(LAS
   ,75),(MCO,74),(SEA,71),...
7 out-degree :: (ATL,151),(ORD,148),(DEN,124),(MSP,112),(DFW,106),(IAH
   ,103),(DTW,99),(PHX,89),(SLC,88),(EWR,85),(SFO,79),(LAX,79),(LAS
   ,75),(MCO,74),(SEA,71),...
8
9 3. Find the top 10 airports with the highest in-degree and out-degree
   values.
10 in-degree :: (ATL,151),(ORD,148),(DEN,124),(MSP,112),(DFW,106),(IAH
   ,103),(DTW,99),(PHX,89),(SLC,88),(EWR,85)
11 out-degree :: (ATL,151),(ORD,148),(DEN,124),(MSP,112),(DFW,106),(IAH
   ,103),(DTW,99),(PHX,89),(SLC,88),(EWR,85)
12
13 4. Find out how many triangles can be formed around Dallas Fort Worth
   International (DFW) airport.
14 Total number of triangles formed around DFW :: 1235
15
16 5. Find out the top 10 airports with the highest page ranks.
17 Top 10 airports with the highest page ranks :: (ATL
   ,10.846702996395441),(ORD,10.100217228379133),(MSP
   ,7.753265350768704),(DEN,7.512767694389999),(IAH
   ,6.458574807471953),(DFW,5.980831089994658),(DTW
   ,5.927876189482903),(SLC,5.816666616332601),(PHX
   ,4.9464141894103575),(SFO,4.629206834210625)
```

18
19 6. Triplets $[A, B, C]$ such that, at least one flight from A to B, and
20 from B to C, but no flight from A to C.
[MCO, SFO, MCO], [PBI, ORD, MIA], [ORF, CLT, SAT], [FLL, TPA, MHT], [IND, DEN, SDF
], [CMH, ATL, RST], [CMH, MSP, BMI], [MKE, LAS, PIT], [OKC, SEA, BOS], [MSN, ORD
, ROA], ...